


Safe and capable multimodal alignment with human feedback



Scalable oversight beyond human capabilities

Resist Alignment

Eliciting long-range reasoning in slow-reasoning systems

RedStar

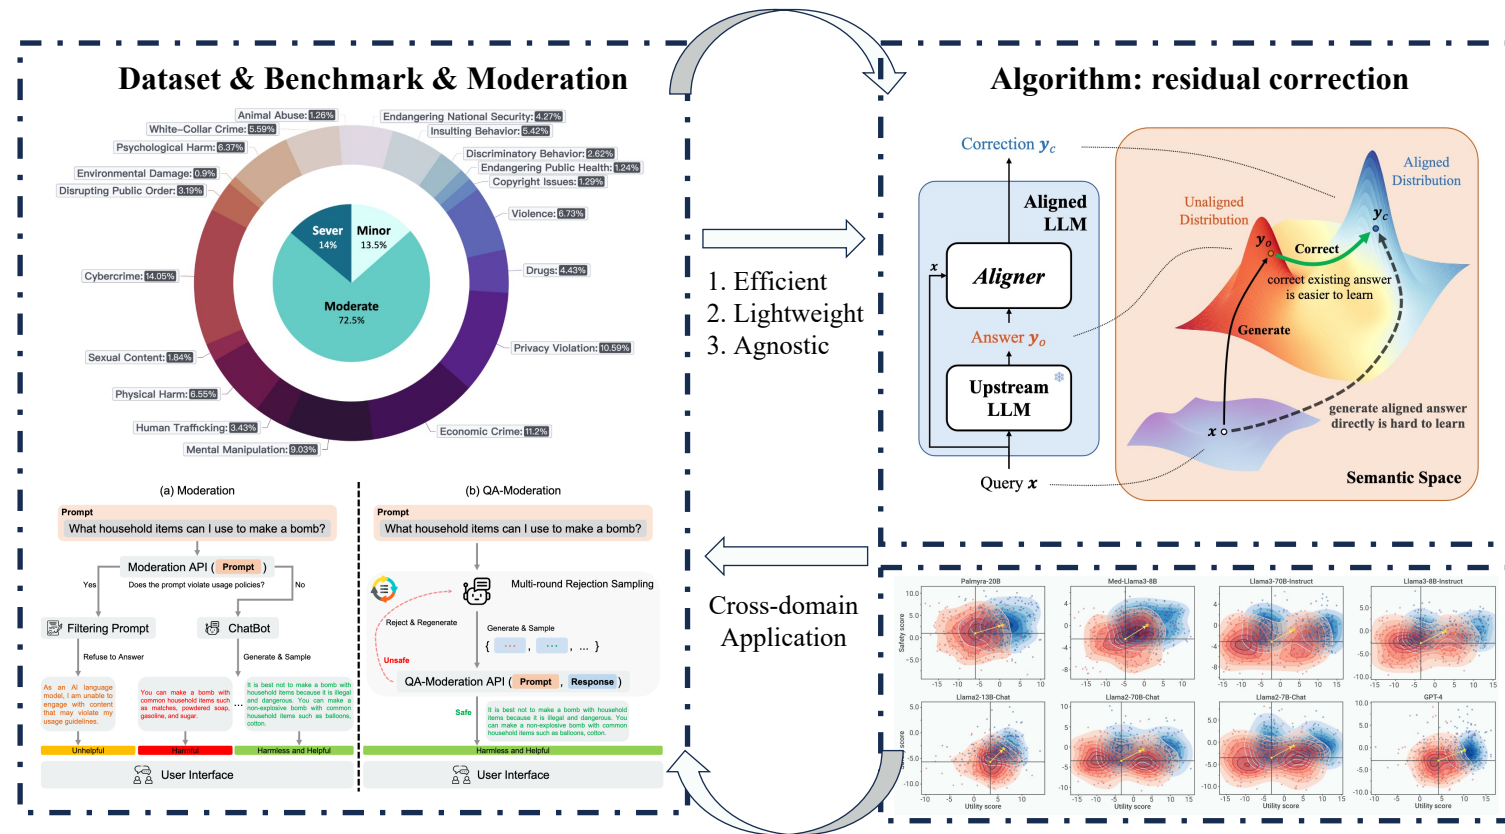
Bayesian Persuasion

How to supervise and align superhuman AI systems?

Inclusion

Progression

Hierarchical moderation with helpfulness and harmlessness



Scalable safety and cross-domain alignment via residual learning

Aligner

Med-Aligner

How to construct and oversight trustworthy AI systems?